

Review Test Submission: ICTPRG554<p style='font-style:italic; font-size:14px;' id='no_display' >If you have any queries regarding this assessment result, comments and suggestions, please contact your instructor/trainer. If we do not hear from you by the end of term, we will assume your acknowledgement and acceptance of this feedback.</p>

Review Test Submission: ICTPRG554

If you have any queries regarding this assessment result, comments and suggestions, please contact your instructor/trainer. If we do not hear from you by the end of term, we will assume your acknowledgement and acceptance of this feedback.

User	SANGAY THINLEY
Course	Cluster - SaaS: Back-End Dev - ICT50220 (Advanced Programming)
Test	ICTPRG554
Started	28/11/24 13:21
Submitted	28/11/24 13:41
Due Date	05/12/24 16:30
Status	Completed
Attempt Outcome	79 out of 100 satisfactory
Time Elapsed	20 minutes out of 1 hour and 30 minutes
Results Displayed	All Answers (Select this one), Submitted Answers (Select this one), Feedback (Select this one), Incorrectly Answered Questions (Select this one)

Question 1

0 out of 7 satisfactory



Scenario: A global logistics company is developing a real-time tracking system that needs to handle diverse data types from multiple sources including GPS trackers, shipping manifests, customs documentation, and customer interactions.

Question: What key advantages does a NoSQL database provide for managing complex, evolving data structures in this logistics scenario?

- Selected Answers:

Schema-free data persistence allows the logistics system to store varied data types without predefined table structures, enabling rapid adaptation to changing business requirements. This flexibility means new types of shipping information can be added without extensive database redesign, supporting the dynamic nature of global logistics data.

Traditional relational databases offer superior data integrity through strict schema enforcement, ensuring that all shipping and tracking data conform to a predefined, uniform structure. This approach provides consistent data validation and prevents inconsistent or incomplete data entry across multiple operational systems.

NoSQL databases support horizontal scalability, allowing the logistics platform to distribute data across multiple servers seamlessly. This approach enables handling increasing volumes of tracking data from global shipping routes, supporting concurrent access from different regional offices and tracking systems without performance bottlenecks.
- Answers:

Schema-free data persistence allows the logistics system to store varied data types without predefined table structures, enabling rapid adaptation to changing business requirements. This flexibility means new types of shipping information can be added without extensive database redesign, supporting the dynamic nature of global logistics data.

Relational databases are more suitable for complex logistics data, as they provide comprehensive support for multi-table joins and complex query operations across shipping, customer, and tracking information. The rigid schema ensures data consistency and simplifies complex reporting requirements.

Traditional relational databases offer superior data integrity through strict schema enforcement, ensuring that all shipping and tracking data conform to a predefined, uniform structure. This approach provides consistent data validation and prevents inconsistent or incomplete data entry across multiple operational systems.

NoSQL databases support horizontal scalability, allowing the logistics platform to distribute data across multiple servers seamlessly. This approach enables handling increasing volumes of tracking data from global shipping routes, supporting concurrent access from different regional offices and tracking systems without performance bottlenecks.

Response

Feedback: **✗ Incorrect**

Hints:

Highlight the key benefit of NoSQL's schema-free, dynamic nature of NoSQL databases in handling evolving data structures.

Question 2

0 out of 7 satisfactory



Scenario: A global content delivery network (CDN) is designing its infrastructure to handle increasing multimedia content storage and delivery across multiple geographic regions.

Question: What are the critical considerations when implementing scaling strategies for a distributed content delivery system?

Select ALL correct statements.

Selected Answers: Horizontal scaling enables the CDN to distribute content across multiple servers and geographic locations, improving content delivery speed and system resilience. This approach allows for dynamic addition of new servers to handle increased traffic and storage requirements without significant system redesign.

Vertical scaling has inherent limitations in distributed systems, as there are physical and economic constraints to upgrading a single server's resources. The approach becomes increasingly expensive and complex as the system requires more computational power and storage capacity.

Hybrid scaling approaches can combine horizontal and vertical scaling strategies, allowing the CDN to optimize performance and cost-effectiveness. This method involves selectively upgrading individual server resources while also distributing load across multiple machines.

Answers: Horizontal scaling enables the CDN to distribute content across multiple servers and geographic locations, improving content delivery speed and system resilience. This approach allows for dynamic addition of new servers to handle increased traffic and storage requirements without significant system redesign.

Vertical scaling has inherent limitations in distributed systems, as there are physical and economic constraints to upgrading a single server's resources. The approach becomes increasingly expensive and complex as the system requires more computational power and storage capacity.

Scaling a distributed system requires comprehensive redesign of the entire application architecture, making it impractical for systems with established infrastructure. The complexity of implementing distributed scaling makes it unsuitable for most real-world applications.

Hybrid scaling approaches can combine horizontal and vertical scaling strategies, allowing the CDN to optimize performance and cost-effectiveness. This method involves selectively upgrading individual server resources while also distributing load across multiple machines.

Response

Feedback: **✗ Incorrect**

Hints:

The question explores the fundamental differences between horizontal and vertical scaling, highlighting their respective strengths, limitations, and appropriate use cases in distributed system design.

The scenario demonstrates that effective scaling strategies require a nuanced understanding of an organization's specific performance, cost, and architectural requirements.

Question 3

7 out of 7 satisfactory



Scenario: A financial services application is developing a high-performance trading analytics system using a NoSQL database to process real-time market data and historical trading records.

Question: What are the critical considerations for optimizing NoSQL database interactions from the programming perspective?

Select ALL correct statements.

Selected Answers: Efficient query design involves minimizing the amount of data transferred between the client and database, using projection and selective retrieval techniques. This approach reduces network overhead and improves overall system performance.

Leveraging database-native aggregation pipelines and server-side query processing can significantly reduce client-side computational overhead. By pushing complex data transformations to the database layer, applications can achieve more efficient data processing.

Answers: NoSQL databases do not support any form of query optimization or index utilization, requiring developers to implement all data filtering and processing logic within the application code. This approach provides maximum flexibility but can lead to performance challenges.

Client-side query optimization requires complete reimplementing of database query logic in the application layer, bypassing the database's native query optimization mechanisms. This ensures maximum control over data retrieval processes.

Efficient query design involves minimizing the amount of data transferred between the client and database, using projection and selective retrieval techniques. This approach reduces network overhead and improves overall system performance.

Leveraging database-native aggregation pipelines and server-side query processing can significantly reduce client-side computational overhead. By pushing complex data transformations to the database layer, applications can achieve more efficient data processing.

Response

Feedback: **✓ Correct**

Question 4

0 out of 7 satisfactory



Scenario: A global social media platform needs to design a partitioning strategy for storing and retrieving billions of user posts across multiple database servers.

Question: What are the key considerations when designing partition keys for large-scale data distribution?

Select ALL correct statements.

- Selected Answers:
- Partition keys should be chosen based on the most common query patterns, considering how data will be accessed and retrieved. For social media posts, this might involve creating compound partition keys that facilitate efficient retrieval of user-specific content.
 - Partition key design should ensure even distribution of data across all database nodes, avoiding hot spots that could lead to performance bottlenecks. A well-designed partition key might combine user ID with timestamp to distribute social media posts evenly across the cluster.
- Answers:
- Partition keys should be chosen based on the most common query patterns, considering how data will be accessed and retrieved. For social media posts, this might involve creating compound partition keys that facilitate efficient retrieval of user-specific content.
 - Random partition key generation is the best approach for distributing social media data, as it eliminates any potential patterns in data access and ensures completely uniform distribution across all nodes.
 - Partition key design should ensure even distribution of data across all database nodes, avoiding hot spots that could lead to performance bottlenecks. A well-designed partition key might combine user ID with timestamp to distribute social media posts evenly across the cluster.
 - The most effective partition key is always the auto-incrementing primary key, as it guarantees sequential data distribution and optimal query performance. This approach ensures that related social media posts are always stored together on the same partition.

Response

Feedback: **✗ Incorrect**

Hints:

Think about data distribution patterns and access requirements.

Consider which options would help achieve balanced data distribution while supporting efficient data retrieval based on common query patterns.

Question 5

7 out of 7 satisfactory



Scenario: An IoT device monitoring system needs to implement efficient data lifecycle management for sensor readings while maintaining system performance.

Question: What considerations are crucial for optimizing TTL performance in a high-throughput IoT environment?

Select ALL correct statements.


- Selected Answers:
- MongoDB's TTL index can be combined with other indexes to optimize query performance while maintaining automatic document expiration. This approach enables efficient data access patterns while ensuring proper lifecycle management.
 - TTL background processes can be configured to run during specific time windows, allowing the system to manage data cleanup during periods of lower activity. This helps maintain optimal system performance during peak operational hours.

Answers: MongoDB's TTL index can be combined with other indexes to optimize query performance while maintaining automatic document expiration. This approach enables efficient data access patterns while ensuring proper lifecycle management.

Implementing TTL requires dedicated servers that continuously scan for expired documents, completely separate from the main database infrastructure. This isolation ensures that TTL operations don't impact regular database performance.

TTL operations must be manually triggered every time data needs to be expired, ensuring complete control over the cleanup process but requiring significant administrative overhead.

TTL background processes can be configured to run during specific time windows, allowing the system to manage data cleanup during periods of lower activity. This helps maintain optimal system performance during peak operational hours.

Response
Feedback:  **Correct**

Question 6

7 out of 7 satisfactory



Scenario: A financial technology platform needs to implement secure authentication for their trading application, ensuring compliance with regulatory requirements while maintaining system performance.

Question: What are the critical factors in selecting and implementing authentication mechanisms?

Select ALL correct statements.

Selected
Answers: Multiple authentication mechanisms can be configured simultaneously, allowing different client applications to use the most appropriate method for their security requirements. This flexibility supports diverse integration scenarios while maintaining security standards.


SCRAM-SHA-256 authentication provides strong security through salted challenge-response mechanisms, making it suitable for protecting sensitive financial data. This method ensures secure credential verification without transmitting passwords in plain text.

Answers: Multiple authentication mechanisms can be configured simultaneously, allowing different client applications to use the most appropriate method for their security requirements. This flexibility supports diverse integration scenarios while maintaining security standards.

Authentication should be disabled during development and testing phases to simplify the development process, then enabled only in production environments. This approach streamlines the development workflow.

All authentication credentials should be stored in environment variables and configuration files that are easily accessible to all development team members, simplifying system maintenance and troubleshooting.

SCRAM-SHA-256 authentication provides strong security through salted challenge-response mechanisms, making it suitable for protecting sensitive financial data. This method ensures secure credential verification without transmitting passwords in plain text.

Response
Feedback:  **Correct**

Question 7

7 out of 7 satisfactory



Scenario: An e-commerce analytics platform needs to implement load balancing for customer behavior data across multiple database partitions.

Question: What are the essential aspects of maintaining balanced data distribution?

Select ALL correct statements.

- Selected Answers:
- Dynamic load balancing mechanisms should monitor partition sizes and automatically redistribute data when imbalances occur. This ensures optimal resource utilization and prevents performance bottlenecks.
 - Load balancing strategies should consider both data volume and access patterns, potentially redistributing hot data across multiple partitions to prevent overload. This approach helps maintain consistent performance.
- Answers:
- Dynamic load balancing mechanisms should monitor partition sizes and automatically redistribute data when imbalances occur. This ensures optimal resource utilization and prevents performance bottlenecks.
 - The system should maintain exactly equal partition sizes at all times, requiring immediate data migration whenever any partition grows larger than others. This ensures perfect balance across the system.
 - Load balancing strategies should consider both data volume and access patterns, potentially redistributing hot data across multiple partitions to prevent overload. This approach helps maintain consistent performance.
 - Load balancing should only be performed during system initialization, with no modifications allowed during operation to ensure system stability.

Response

Feedback:  **Correct**

Question 8

7 out of 7 satisfactory



Scenario: A financial reporting system needs to implement comprehensive testing for its MongoDB-based data aggregation pipelines.

Question: What are the key considerations for developing an effective NoSQL testing strategy?

Select ALL correct statements.

- Selected Answers:
- Integration tests should verify data consistency across different database operations, including testing of atomic transactions and data integrity constraints. This ensures reliable system behavior under various conditions.
 - Test environments should include representative data samples that match production data patterns and volumes. This enables realistic testing of query performance and data manipulation operations.
- Answers:
- Unit tests should focus exclusively on application code, ignoring database interactions as they are too complex to test effectively. This simplifies the testing process and reduces maintenance overhead.
 - Integration tests should verify data consistency across different database operations, including testing of atomic transactions and data integrity constraints. This ensures reliable system behavior under various conditions.
 - Test environments should include representative data samples that match production data patterns and volumes. This enables realistic testing of query performance and data manipulation operations.
 - Testing should only be performed on production databases to ensure accurate results, as test environments cannot properly replicate real-world conditions.

Response

Feedback:  **Correct**

Question 9

7 out of 7 satisfactory



Scenario: A time-series analytics platform needs to implement efficient data retrieval for historical sensor readings stored in a NoSQL database.

Question: What are the key considerations when designing sort keys for time-series data?

Select ALL correct statements.

- Selected Answers:
- Compound sort keys can combine timestamp data with additional attributes like sensor ID or location, enabling efficient querying across multiple dimensions. This supports complex analytics requirements while maintaining performance.
 - Sort keys can be designed to support range-based queries, allowing efficient retrieval of sensor readings within specific time periods. This enables fast access to historical data while maintaining optimal storage organization.
- Answers:
- Creating sort keys should always use random values to ensure even data distribution, regardless of the natural ordering of the data. This approach maximizes storage efficiency across all partitions.
 - Compound sort keys can combine timestamp data with additional attributes like sensor ID or location, enabling efficient querying across multiple dimensions. This supports complex analytics requirements while maintaining performance.
 - Sort keys can be designed to support range-based queries, allowing efficient retrieval of sensor readings within specific time periods. This enables fast access to historical data while maintaining optimal storage organization.
 - Sort keys must be updated regularly to maintain system performance, requiring complete data reorganization on a weekly basis. This ensures optimal query performance over time.

Response

Feedback: **Correct**

Question 10

7 out of 7 satisfactory



Scenario: A financial services platform needs to implement secure authentication mechanisms for their distributed NoSQL database system.

Question: What are the essential aspects of selecting and implementing authentication protocols?

Select ALL correct statements.

- Selected Answers:
- SCRAM-SHA-256 authentication provides strong security through challenge-response mechanisms and salted password hashing. This protects against various attack vectors while ensuring secure credential verification.
 - X.509 certificate authentication can be implemented for client-server and cluster member authentication, providing strong identity verification. This enables secure communication between all system components.
- Answers:
- SCRAM-SHA-256 authentication provides strong security through challenge-response mechanisms and salted password hashing. This protects against various attack vectors while ensuring secure credential verification.
 - Authentication mechanisms should be identical for all users and services, regardless of their security requirements or access patterns. This ensures consistent system behavior.
 - X.509 certificate authentication can be implemented for client-server and cluster member authentication, providing strong identity verification. This enables secure communication between all system components.
 - Basic authentication using plaintext passwords is sufficient for all database connections, as transport encryption provides adequate security. This simplifies the authentication process.

Response
Feedback:  **Correct**

Question 11

7 out of 7 satisfactory



Scenario: A fraud detection system needs to implement pattern recognition using graph database capabilities.

Question: What are the key advantages of using graph databases for pattern analysis?

Select ALL correct statements.

Selected Answers: Graph databases excel at identifying complex relationships and patterns in connected data, making them ideal for detecting suspicious transaction networks and fraud rings. Their native support for relationship traversal enables efficient pattern matching.

Real-time graph traversal capabilities enable immediate detection of suspicious patterns as new transactions occur, supporting proactive fraud prevention. The database can efficiently explore relationship networks to identify potential fraud indicators.

Answers: Graph databases excel at identifying complex relationships and patterns in connected data, making them ideal for detecting suspicious transaction networks and fraud rings. Their native support for relationship traversal enables efficient pattern matching.

Graph databases can only store simple, linear relationships and are not suitable for complex data analysis scenarios. They should be avoided for any serious analytical applications.

Real-time graph traversal capabilities enable immediate detection of suspicious patterns as new transactions occur, supporting proactive fraud prevention. The database can efficiently explore relationship networks to identify potential fraud indicators.

Graph databases require all data to be stored in predefined, rigid structures that cannot be modified once established. This makes them inflexible for evolving fraud detection needs.

Response
Feedback:  **Correct**

Question 12

7 out of 7 satisfactory



Scenario: A global event management platform needs to implement robust date-time handling for events scheduled across multiple time zones.

Question: What considerations are crucial for managing date-time data in a NoSQL database?

Select ALL correct statements.

Selected Answers: DateTime values should be stored in a standardized format (like UTC) with time zone information preserved separately. This enables accurate time zone conversion and consistent event scheduling across regions.

Date-time indexing strategies should support efficient range queries for event scheduling and calendar views. This allows quick retrieval of events within specific time periods.

Answers: DateTime values should be stored in a standardized format (like UTC) with time zone information preserved separately. This enables accurate time zone conversion and consistent event scheduling across regions.

All dates should be stored as simple string values without any timezone information, relying on application logic to handle time zone conversions. This simplifies database operations.

Temporal data should always be stored as local time values without any timezone context, as this provides the most natural representation for users.

Date-time indexing strategies should support efficient range queries for event scheduling and calendar views. This allows quick retrieval of events within specific time periods.

Response

Feedback:  **Correct**

Question 13

8 out of 8 satisfactory



Scenario: A smart city initiative is developing an IoT analytics platform to monitor and analyze data from millions of sensors across urban infrastructure, including traffic systems, environmental monitors, and utility meters. The platform needs to process real-time data streams, maintain historical analytics, ensure data security, and provide efficient query capabilities for both real-time monitoring and long-term trend analysis.

Question: Which combination of architectural decisions and implementation strategies would create the most effective and scalable solution for this IoT platform?

Select the FIVE most appropriate options.

Selected Answers: Implement a hybrid storage solution using MongoDB time-series collections for real-time sensor data, with automated archival processes moving historical data to optimized collections based on access patterns.

Deploy a graph database component to analyze relationships between different urban systems, enabling complex pattern recognition for infrastructure interdependencies and cascade effect analysis.

Create a comprehensive testing framework that simulates various failure scenarios, including network partitions and hardware failures, while validating data consistency and recovery procedures across distributed nodes.

Design a partitioning strategy that combines sensor type and geographic location in the partition key, enabling efficient data distribution and queries while maintaining data locality for regional analytics.

Implement a multi-layer security strategy combining certificate-based authentication, field-level encryption for sensitive measurements, and role-based access control for different stakeholder groups.

Answers: Implement a hybrid storage solution using MongoDB time-series collections for real-time sensor data, with automated archival processes moving historical data to optimized collections based on access patterns.

Deploy a graph database component to analyze relationships between different urban systems, enabling complex pattern recognition for infrastructure interdependencies and cascade effect analysis.

Create a comprehensive testing framework that simulates various failure scenarios, including network partitions and hardware failures, while validating data consistency and recovery procedures across distributed nodes.

Store all sensor data in a single collection without considering data lifecycle management or access patterns, maximizing storage utilization.

Disable data validation and integrity checks to improve ingestion performance, assuming all sensor data is accurate and properly formatted.

Design a partitioning strategy that combines sensor type and geographic location in the partition key, enabling efficient data distribution and queries while maintaining data locality for regional analytics.

Implement a multi-layer security strategy combining certificate-based authentication, field-level encryption for sensitive measurements, and role-based access control for different stakeholder groups.

Rely on client-side sorting and filtering for all analytics queries, minimizing database processing requirements.

Use basic authentication mechanisms without encryption for internal system components, assuming network-level security is sufficient.

Implement static data partitioning without considering data growth or access patterns, focusing on immediate deployment requirements. Implement static data partitioning without considering data growth or access patterns, focusing on immediate deployment requirements.

Response

Feedback:  **Correct**

Question 14

8 out of 8 satisfactory



Scenario: A global retail organization is planning to migrate their existing enterprise resource planning (ERP) system from a traditional relational database to a NoSQL solution. The system handles millions of transactions daily, stores complex product hierarchies, manages user authentication across multiple regions, and needs to maintain historical data with automatic cleanup processes. The organization needs to ensure high availability, data security, and optimal performance across all regions.

Question: Considering the comprehensive requirements for this migration, which combination of approaches would provide the most robust and efficient solution?

Select the FIVE most appropriate options.

- | | |
|-------------------|---|
| Selected Answers: | <p>Define composite indexes combining product hierarchy and regional identifiers, supporting efficient query patterns for both local and global inventory management.</p> <p>Design a hybrid scaling strategy that combines horizontal scaling for transaction processing with vertical scaling for analytics workloads, optimizing resource utilization across different workload types.</p> <p>Configure SCRAM-SHA-256 authentication with X.509 certificates for inter-cluster communication, implementing role-based access control with field-level encryption for sensitive data.</p> <p>Utilize TTL indexes for transaction history management, automatically expiring old records based on regulatory requirements while maintaining recent data for active analysis.</p> <p>Implement a document-based NoSQL store using MongoDB's native sharding capabilities to distribute data across geographic regions, ensuring data locality and reduced latency for regional access patterns.</p> |
| Answers: | <p>Define composite indexes combining product hierarchy and regional identifiers, supporting efficient query patterns for both local and global inventory management.</p> <p>Design a hybrid scaling strategy that combines horizontal scaling for transaction processing with vertical scaling for analytics workloads, optimizing resource utilization across different workload types.</p> <p>Employ sort keys based on timestamp and region identifiers, enabling efficient range-based queries for regional performance analysis and reporting.</p> <p>Configure SCRAM-SHA-256 authentication with X.509 certificates for inter-cluster communication, implementing role-based access control with field-level encryption for sensitive data.</p> <p>Utilize TTL indexes for transaction history management, automatically expiring old records based on regulatory requirements while maintaining recent data for active analysis.</p> |

Use random partition keys for data distribution, ignoring natural data access patterns and regional requirements.

Implement a document-based NoSQL store using MongoDB's native sharding capabilities to distribute data across geographic regions, ensuring data locality and reduced latency for regional access patterns.

Implement basic authentication using plain text passwords, assuming transport layer encryption provides sufficient security for all database connections.

Store all data in a single, centralized database instance to simplify management and reduce operational complexity, relying on application-level caching for performance optimization.

Maintain separate authentication mechanisms for each region, with no centralized access control or consolidated security policies.

Response

Feedback:  **Correct**

Thursday, 28 November 2024 13:41:23 o'clock AWST

← OK